

# Computer Operating Systems Homework 3

Hakan Duran 150200091

May 2024

Your report should convey the problem and the implementation details of your code.

## 1 Introduction & Problem

In this homework, a single-threaded web server is given and we were asked to make this server multi-threaded. By using thread APIs, i developed skills on multi-threading and observed how threads add speed to the process, especially for web servers.

## 2 Implementation details

### 2.1 server.c

In my code, i didn't add the code for extra part, which is scheduling policies (FIFO etc).

As an introduction, i need to introduce global variables and initialization for threads and semaphores:

```
1  #include <pthread.h>
2  #include <semaphore.h>
3  #include "blg312e.h"
4  #include "request.h"
5
6  sem_t empty; // Semaphore to track empty slots in the buffer
7  sem_t full; // Semaphore to track filled slots in the buffer
8  sem_t mutex; // Mutex for accessing shared buffer
9  int *buffer; // Shared buffer to hold connection descriptors
10 int num_buffers = 0; // Maximum number of buffers
11 int in = 0; // Index for adding items to the buffer
12 int out = 0; // Index for removing items from the buffer
13
14 int main(int argc, char *argv[])
15 {
16     int listenfd, connfd, port, clientlen;
17     struct sockaddr_in clientaddr;
18     int num_threads;
19
20     getargs(&port, &num_threads, &num_buffers, argc, argv);
21
22     buffer = malloc(num_buffers * sizeof(int));
23     if (buffer == NULL) {
24         fprintf(stderr, "Failed to allocate memory for buffer\n");
25         exit(1);
26     }
27
28     // Semaphores initialized here
29     sem_init(&empty, 0, num_buffers); // It will make thread sleep when buffer is full
30     sem_init(&full, 0, 0); // It will make thread sleep when buffer is empty
31     sem_init(&mutex, 0, 1); // Locks for producer-consumer relationship
32
```

```

33 // Threads created here
34 pthread_t tid;
35 for (int i = 0; i < num_threads; i++) {
36     pthread_create(&tid, NULL, worker, NULL);
37 }
38 ...
39 }

```

There are 3 semaphores here called as empty, full and mutex. Empty and full semaphores are for producer-consumer relationship between worker and master thread. Mutex is used for critical sections.

Buffer pointer points to buffer array, which created in line 22. This buffer will consist of file descriptors which will be reached by worker threads.

Semaphore initialization has been made carefully as it is seen in line 29. Semaphore "empty" is initialized by the value of num\_buffers. It is decremented when new request comes in. If buffer became full, it will make master thread sleep. Semaphore "full" is initialized by the value of 0. It will make worker threads sleep when buffer is empty. Since semaphore "mutex" is for lock, it is initialized by 1.

In line 35, it can be seen that threads are being created by using the function of pthread\_create(). The number of threads is determined by the user.

```

1 void *worker(void *arg) {
2     int connfd;
3     while (1) {
4         sem_wait(&full); // Wait if buffer is empty
5         sem_wait(&mutex);
6         connfd = buffer[out]; // Get connection descriptor from buffer
7         out = (out + 1) % num_buffers; // Update index
8         sem_post(&mutex);
9         sem_post(&empty); // Signal that there's an empty slot in the buffer
10
11         // Handle the request
12         requestHandle(connfd);
13         Close(connfd);
14     }
15     return NULL;
16 }
17
18 int main(int argc, char *argv[])
19 {
20     ...
21     listenfd = Open_listenfd(port);
22     while (1) {
23         clientlen = sizeof(clientaddr);
24         connfd = Accept(listenfd, (SA *)&clientaddr, (socklen_t *)&clientlen);
25
26         sem_wait(&empty); // Wait if buffer is full
27         sem_wait(&mutex);
28         buffer[in] = connfd; // Put connection descriptor in buffer
29         in = (in + 1) % num_buffers; // Update index
30         sem_post(&mutex);
31         sem_post(&full); // Signal that there's a filled slot in the buffer
32     }
33 }

```

In the above code, worker and master threads can be seen. Connection is accepted and connection file descriptor (connfd) put in the buffer while request is handled by worker threads.

In the master thread, after getting connection fd (connfd), the first function we see is sem\_wait(&empty), which means master will wait for buffer to empty. It only happens in line 9, sem\_post(&empty), which signals there is empty slot in the buffer.

After first semaphore, we came to line 27, sem\_wait(&mutex). Mutual exclusion prevents threads from interfering in each other's jobs. File descriptor is added to buffer and counter variable "in" updated.

Semaphore function sem\_post(&mutex) shows that it is end of critical section. At line 31, sem\_post(&full) function increases the semaphore "full" by one, signals to the worker threads which waits at line 4, sem\_wait(&full).

Worker thread is not much different than master thread. It waits for a signal from master if buffer is empty. In critical section, connfd has acquired and counter variable "out" updated. In the last, request handled.

## 2.2 client.c

In this part, i am extracting requested files from a file. This file is consists of requested files and seperated by newline. The name is given from console.

Usage:

```
./client <host> <port> <file>
```

Example usage:

```
./client localhost 2000 files.txt
```

Inside of files.txt:

```
1 /output.cgi
2 /favicon.ico
3 /output.cgi
4 /
5 /output.cgi
```

Here is the part of main function:

```
1 char line[MAX_FILENAME_LENGTH];
2
3 // In here, each line of file is assigned to "line" variable
4 // Connection established with Open_clientfd
5 // Parameters are determined and thread created
6
7 while (fgets(line, sizeof(line), file) != NULL) {
8
9     line[strcspn(line, "\n")] = '\0';
10
11     clientfd = Open_clientfd(host, port);
12
13     ThreadParams *params = malloc(sizeof(ThreadParams));
14     params->arg1 = clientfd;
15     params->arg2 = strdup(line);
16
17     pthread_t tid;
18     pthread_create(&tid, NULL, clientSend, (void *)params);
19
20 }
```

Each line extracted from file and clientfd are parameters for created thread. Below is the clientSend function:

```
1 typedef struct { // Struct for clientSend function's parameters
2     int arg1;
3     char* arg2;
4 } ThreadParams;
5
6 void clientSend(void *args)
7 {
8     // Parameters are assigned to values here
9     ThreadParams *params = (ThreadParams *)args;
10     int fd = params->arg1;
11     char* filename = params->arg2;
12
13     char buf1[MAXLINE];
14     char hostname[MAXLINE];
15
16     .
17     .
18     .
19
20     while (n > 0) {
```

```

21     printf("%s", buf);
22     n = Rio_readlineb(&rio, buf, MAXBUF);
23 }
24
25 // Allocated params value is freed here
26 free(params);
27
28 return NULL;
29 }

```

In first 4 lines, a struct created for sending parameters to threads. Line 8-11 is for assigning parameter values to the inner variables. In line 26, params is freed.

### 3 Example

In order to test my code, i've requested files in files.txt. I also changed spinfor variable from 5 to 10 in order to have much more scalable measurement in output.c. Here are my server and clients:

1. ./server 2000 5 10
2. ./client localhost 2000 files.txt

Inside of files.txt:

```

1 /output.cgi
2 /favicon.ico
3 /output.cgi
4 /
5 /output.cgi

```

Before implementing of multi-threading, it would take roughly 30 seconds to handle all requests. After implementing multi-threading, it started to take roughly 10 seconds. It is because different threads have a chance to deal with different requests anymore.

Here is a pcap file from wireshark to show how requesting 5 file which consists of 3 output.cgi is only takes 10 seconds:

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	127.0.0.1	127.0.0.1	TCP	74	45480 → 2000 [SYN] Seq=0 Win=65495 Len=0 MSS=65495 SACK_PERM=1 TSval=4195739879 TSecr=0 WS=128
2	0.000000052	127.0.0.1	127.0.0.1	TCP	74	2000 → 45480 [RST, ACK] Seq=0 Ack=1 Win=0 Len=0 MSS=65495 SACK_PERM=1 TSval=4195739879 TSecr=4195739879 WS=128
3	0.000000138	127.0.0.1	127.0.0.1	TCP	66	45480 → 2000 [ACK] Seq=1 Ack=1 Win=65536 Len=0 TSval=4195739879 TSecr=4195739879
4	0.000000230	127.0.0.1	127.0.0.1	TCP	74	45480 → 2000 [SYN] Seq=0 Win=65495 Len=0 MSS=65495 SACK_PERM=1 TSval=4195739879 TSecr=0 WS=128
5	0.000000711	127.0.0.1	127.0.0.1	TCP	74	2000 → 45480 [ACK] Seq=0 Ack=1 Win=65536 Len=0 MSS=65495 SACK_PERM=1 TSval=4195739879 TSecr=4195739879 WS=128
6	0.000000864	127.0.0.1	127.0.0.1	HTTP	115	GET /output.cgi HTTP/1.1
7	0.000000912	127.0.0.1	127.0.0.1	TCP	66	45480 → 2000 [ACK] Seq=1 Ack=1 Win=65536 Len=0 TSval=4195739879 TSecr=4195739879
8	0.000001375	127.0.0.1	127.0.0.1	TCP	66	2000 → 45480 [ACK] Seq=1 Ack=50 Win=65536 Len=0 TSval=4195739879 TSecr=4195739879
9	0.000001550	127.0.0.1	127.0.0.1	HTTP	116	GET /favicon.ico HTTP/1.1
10	0.000001692	127.0.0.1	127.0.0.1	TCP	66	2000 → 45480 [ACK] Seq=1 Ack=51 Win=65536 Len=0 TSval=4195739879 TSecr=4195739879
11	0.000001798	127.0.0.1	127.0.0.1	TCP	74	45480 → 2000 [SYN] Seq=0 Win=65495 Len=0 MSS=65495 SACK_PERM=1 TSval=4195739879 TSecr=0 WS=128
12	0.000001868	127.0.0.1	127.0.0.1	TCP	74	2000 → 45480 [ACK] Seq=0 Ack=1 Win=65536 Len=0 MSS=65495 SACK_PERM=1 TSval=4195739879 TSecr=4195739879 WS=128
13	0.000001951	127.0.0.1	127.0.0.1	TCP	66	45480 → 2000 [ACK] Seq=1 Ack=1 Win=65536 Len=0 TSval=4195739879 TSecr=4195739879
14	0.000002044	127.0.0.1	127.0.0.1	TCP	111	2000 → 45480 [PSH, ACK] Seq=1 Ack=50 Win=65536 Len=45 TSval=4195739879 TSecr=4195739879 [TCP segment of a reassembled PDU]
15	0.000002093	127.0.0.1	127.0.0.1	TCP	66	45480 → 2000 [ACK] Seq=50 Ack=46 Win=65536 Len=0 TSval=4195739879 TSecr=4195739879
16	0.000001958	127.0.0.1	127.0.0.1	TCP	161	2000 → 45480 [PSH, ACK] Seq=1 Ack=51 Win=65536 Len=95 TSval=4195739879 TSecr=4195739879 [TCP segment of a reassembled PDU]
17	0.000002147	127.0.0.1	127.0.0.1	TCP	66	45480 → 2000 [ACK] Seq=51 Ack=56 Win=65536 Len=0 TSval=4195739879 TSecr=4195739879
18	0.000003891	127.0.0.1	127.0.0.1	TCP	74	45480 → 2000 [SYN] Seq=0 Win=65495 Len=0 MSS=65495 SACK_PERM=1 TSval=4195739880 TSecr=0 WS=128
19	0.000003974	127.0.0.1	127.0.0.1	TCP	74	2000 → 45480 [RST, ACK] Seq=0 Ack=1 Win=0 Len=0 MSS=65495 SACK_PERM=1 TSval=4195739880 TSecr=4195739880 WS=128
20	0.000004100	127.0.0.1	127.0.0.1	TCP	66	45480 → 2000 [ACK] Seq=1 Ack=1 Win=65536 Len=0 TSval=4195739880 TSecr=4195739880 WS=128
21	0.000004063	127.0.0.1	127.0.0.1	TCP	74	45480 → 2000 [SYN] Seq=0 Win=65495 Len=0 MSS=65495 SACK_PERM=1 TSval=4195739880 TSecr=0 WS=128
22	0.000004054	127.0.0.1	127.0.0.1	TCP	74	2000 → 45480 [ACK] Seq=0 Ack=1 Win=65536 Len=0 MSS=65495 SACK_PERM=1 TSval=4195739880 TSecr=4195739880 WS=128
23	0.000005066	127.0.0.1	127.0.0.1	HTTP	115	GET /output.cgi HTTP/1.1
24	0.000005056	127.0.0.1	127.0.0.1	TCP	66	45480 → 2000 [ACK] Seq=1 Ack=1 Win=65536 Len=0 TSval=4195739880 TSecr=4195739880
25	0.000005182	127.0.0.1	127.0.0.1	TCP	66	2000 → 45480 [ACK] Seq=1 Ack=50 Win=65536 Len=0 TSval=4195739880 TSecr=4195739880
26	0.000005068	127.0.0.1	127.0.0.1	HTTP	1472	HTTP/1.0 200 OK (text/plain)
27	0.001042170	127.0.0.1	127.0.0.1	TCP	66	45480 → 2000 [ACK] Seq=51 Ack=1502 Win=64512 Len=0 TSval=4195739880 TSecr=4195739880
28	0.001052006	127.0.0.1	127.0.0.1	HTTP	105	GET / HTTP/1.1
29	0.001052006	127.0.0.1	127.0.0.1	TCP	66	2000 → 45480 [ACK] Seq=1 Ack=40 Win=65536 Len=0 TSval=4195739880 TSecr=4195739880
30	0.001052045	127.0.0.1	127.0.0.1	TCP	111	2000 → 45480 [PSH, ACK] Seq=1 Ack=50 Win=65536 Len=45 TSval=4195739880 TSecr=4195739880 [TCP segment of a reassembled PDU]
31	0.001052041	127.0.0.1	127.0.0.1	TCP	66	45480 → 2000 [ACK] Seq=50 Ack=46 Win=65536 Len=0 TSval=4195739880 TSecr=4195739880
32	0.001052099	127.0.0.1	127.0.0.1	TCP	159	2000 → 45480 [PSH, ACK] Seq=1 Ack=40 Win=65536 Len=93 TSval=4195739881 TSecr=4195739880 [TCP segment of a reassembled PDU]
33	0.001051704	127.0.0.1	127.0.0.1	TCP	66	45480 → 2000 [ACK] Seq=40 Ack=94 Win=65536 Len=0 TSval=4195739881 TSecr=4195739881
34	0.001054508	127.0.0.1	127.0.0.1	HTTP	313	HTTP/1.0 200 OK (text/html)
35	0.001052074	127.0.0.1	127.0.0.1	TCP	66	45480 → 2000 [ACK] Seq=40 Ack=341 Win=65408 Len=0 TSval=4195739881 TSecr=4195739881
36	0.002111148	127.0.0.1	127.0.0.1	HTTP	115	GET /output.cgi HTTP/1.1
37	0.002130537	127.0.0.1	127.0.0.1	TCP	66	2000 → 45480 [ACK] Seq=1 Ack=50 Win=65536 Len=0 TSval=4195739881 TSecr=4195739881
38	0.002040164	127.0.0.1	127.0.0.1	TCP	111	2000 → 45480 [PSH, ACK] Seq=1 Ack=50 Win=65536 Len=45 TSval=4195739881 TSecr=4195739881 [TCP segment of a reassembled PDU]
39	0.002193325	127.0.0.1	127.0.0.1	TCP	66	45480 → 2000 [ACK] Seq=50 Ack=46 Win=65536 Len=0 TSval=4195739881 TSecr=4195739881
40	0.002193956	127.0.0.1	127.0.0.1	HTTP	238	HTTP/1.0 200 OK (text/html)
41	0.007783993	127.0.0.1	127.0.0.1	TCP	66	45480 → 2000 [ACK] Seq=50 Ack=218 Win=65408 Len=0 TSval=4195749887 TSecr=4195749887
42	0.007813814	127.0.0.1	127.0.0.1	TCP	66	2000 → 45480 [FIN, ACK] Seq=341 Ack=40 Win=65536 Len=0 TSval=4195749887 TSecr=4195739881
43	0.008186399	127.0.0.1	127.0.0.1	TCP	238	HTTP/1.0 200 OK (text/html)
44	0.008195987	127.0.0.1	127.0.0.1	TCP	66	45480 → 2000 [ACK] Seq=50 Ack=218 Win=65408 Len=0 TSval=4195749887 TSecr=4195749887
45	0.008245449	127.0.0.1	127.0.0.1	HTTP	238	HTTP/1.0 200 OK (text/html)
46	0.008264527	127.0.0.1	127.0.0.1	TCP	66	45480 → 2000 [ACK] Seq=50 Ack=218 Win=65408 Len=0 TSval=4195749887 TSecr=4195749887
47	0.008395327	127.0.0.1	127.0.0.1	TCP	66	2000 → 45480 [FIN, ACK] Seq=1502 Ack=51 Win=65536 Len=0 TSval=4195749887 TSecr=4195739880
48	0.008478528	127.0.0.1	127.0.0.1	TCP	66	2000 → 45480 [FIN, ACK] Seq=218 Ack=50 Win=65536 Len=0 TSval=4195749887 TSecr=4195749887
49	0.008494585	127.0.0.1	127.0.0.1	TCP	66	2000 → 45480 [FIN, ACK] Seq=218 Ack=50 Win=65536 Len=0 TSval=4195749887 TSecr=4195749887

Figure 1: PCAP file of whole request-response event

After 40. packet, it can be seen that requested output.cgi files has arrived. The arrival of packets containing the FIN flag after the 47th packet proves that it only takes 10 seconds to request 5 files. Since they are threads, small files are coming primarily even though they are not the one which requested firstly.